

Extensions of the Methods for Understanding Neural Network Generalization

Thomas Walker

CONTENTS

I	Introduction	2
1	Components of Neural Network Generalization	2
II	Exploring Stiffness	3
2	Stiffness on Neighbourhoods	3
3	Stiffness on Layers	3
4	Stiffness and Activation	4
III	Appendix	6

AFFILIATION Imperial College London

CORRESPONDENCE thomas.walker21@imperial.ac.uk

DATE July 2023

Part I

Introduction

In this report, I intend to extend some of the work investigated in 'A Survey of the Methods for Deriving Bounds on the Generalization Error of Deep Neural Networks'.

1 COMPONENTS OF NEURAL NETWORK GENERALIZATION

There are three main components of neural networks that are used to understand their generalization capacities. These components are the training data, the training algorithm and the architecture. Some approaches deal exclusively with one of these components while other approaches aim to capture multiple components. For example, information-theoretic approaches focus on the training data and how information is transferred to the representation of the network. On the other hand, Bayesian machine learning is a framework for investigating the learning algorithm. Forming data-inspired priors was an extension of this framework to incorporate training data to get tighter bounds on the generalization. Therefore, to get a comprehensive picture of the neural network generalization we ought to develop ideas that incorporate each of these three components.

Part II

Exploring Stiffness

Recall that stiffness [1] investigates the quantity

$$\bar{g} = \nabla_W \mathcal{L}(f_W(X), y),$$

for a functional approximation f , parameterized by a trainable parameter W . Where $\mathcal{L}(f_W(X), y)$ is the loss function.

Definition 1.1 For two data points (X_1, y_1) and (X_2, y_2) define the sign stiffness to be

$$S_{\text{sign}}((X_1, y_1), (X_2, y_2); f) = \mathbb{E}(\text{sign}(\bar{g}_1 \cdot \bar{g}_2)).$$

Definition 1.2 For two data points (X_1, y_1) and (X_2, y_2) define the cosine stiffness to be

$$S_{\text{cos}}((X_1, y_1), (X_2, y_2); f) = \mathbb{E}(\cos(\bar{g}_1 \cdot \bar{g}_2)),$$

where

$$\cos(\bar{g}_1 \cdot \bar{g}_2) = \frac{\bar{g}_1 \cdot \bar{g}_2}{|\bar{g}_1| |\bar{g}_2|}.$$

Therefore, stiffness captures information from the training data as the gradient of the loss function is evaluated at data points. It also captures information about the learning algorithm as the gradient of the loss function with respect to the parameters determines the update step of stochastic gradient descent. In the subsequent section, I look to investigate how stiffness can be related to the architecture of the neural network.

2 STIFFNESS ON NEIGHBOURHOODS

For a data point, we can investigate the stiffness in relation to the examples within a neighbourhood defined by some metric. For a clustering task, we can consider Euclidean distance to identify the boundaries of a cluster. Refer to Figure 1 for a visualization of this for a two-dimensional example.

3 STIFFNESS ON LAYERS

Instead of representing the parameters as a singular vector, we can decompose them to reflect the architecture of the network. In the following we conduct similar analyses but with the gradient taken with respect to a subset of the parameters to identify how stiffness evolves through the layers of a network.

In the following, we consider the stiffness of data points in relation to a balanced sample of data points from the categories of the dataset. In Figure 2 we compare how the stiffness varies through the layers of a trained and untrained network.

Then, in Figure 3 we look at how stiffness varies through the layers for networks of different architectures.

4 STIFFNESS AND ACTIVATION

In the original experiment, ReLU activation functions connected the layers of the networks. We can instead use tanh and LeakyReLU activation and observe how this affects the stiffness through the layers. My intention here is to understand whether the conclusions of [2] are present when neural networks are investigated from this perspective. Figure 4 shows the results for tanh activation, whereas, Figure 5 shows the results for networks with LeakyReLU activation.

REFERENCES

- [1] Stanislav Fort, Pawel Krzysztof Nowak, and Srini Narayanan. “Stiffness: A New Perspective on Generalization in Neural Networks”. In: *CoRR* (2019).
- [2] Gregory Naitzat, Andrey Zhitnikov, and Lek-Heng Lim. “Topology of deep neural networks”. In: *CoRR* (2020).

Part III

Appendix




Figure 1: The stiffness of points within a neighbourhood can be used to identify the learned boundary of a classification problem.

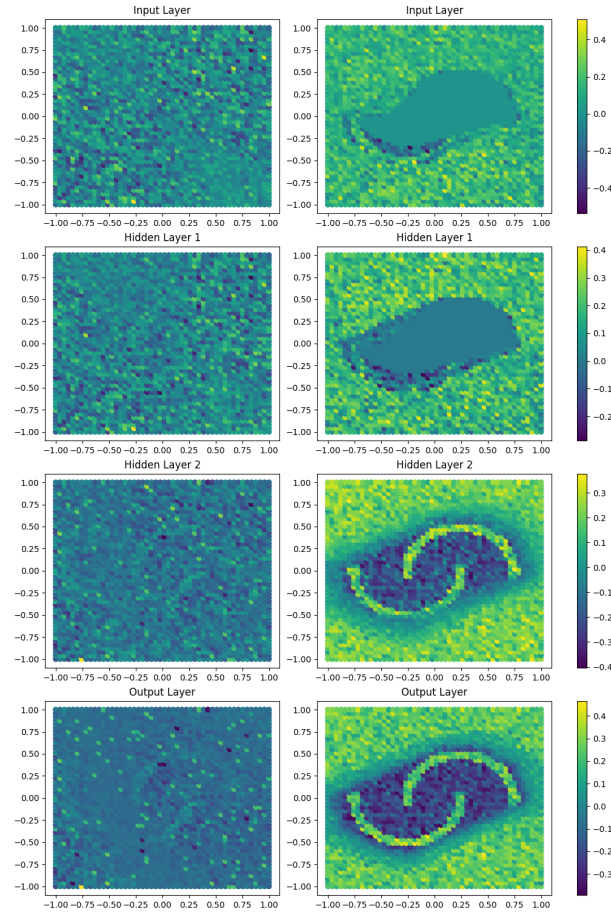


Figure 2: Average stiffness between an example and a random sample from the training data containing samples from both categories.

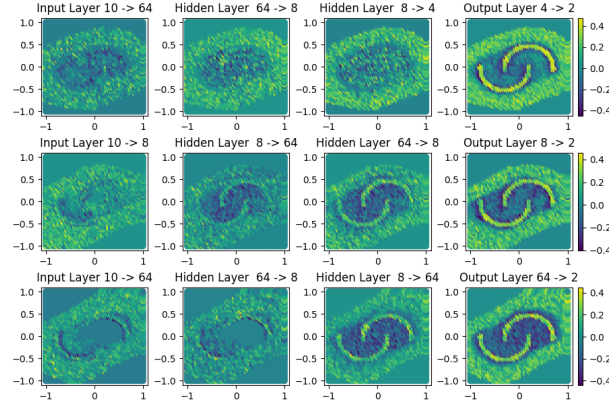


Figure 3: This shows how the stiffness of examples changes through the layers of networks with different numbers of hidden units.

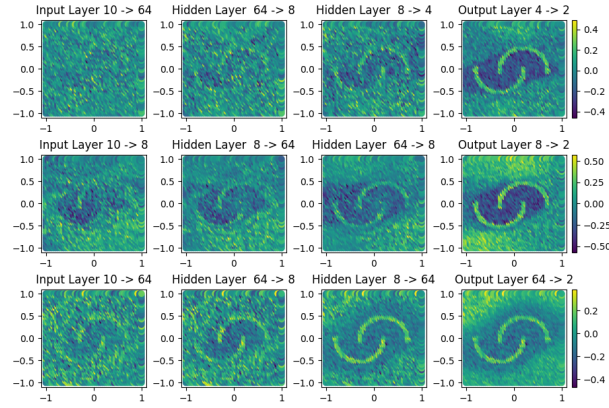


Figure 4: Stiffness through network layers connected by tanh activations

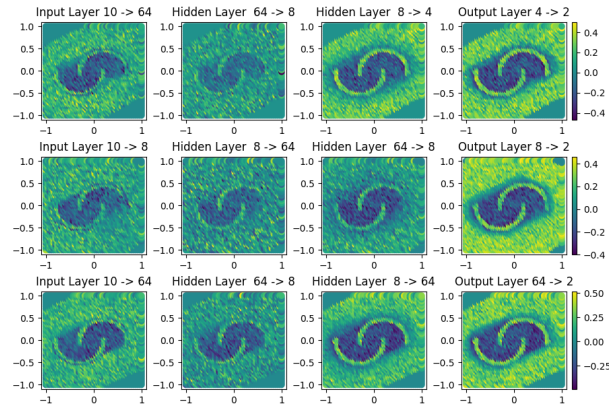


Figure 5: Stiffness through network layers connected by tanh activations